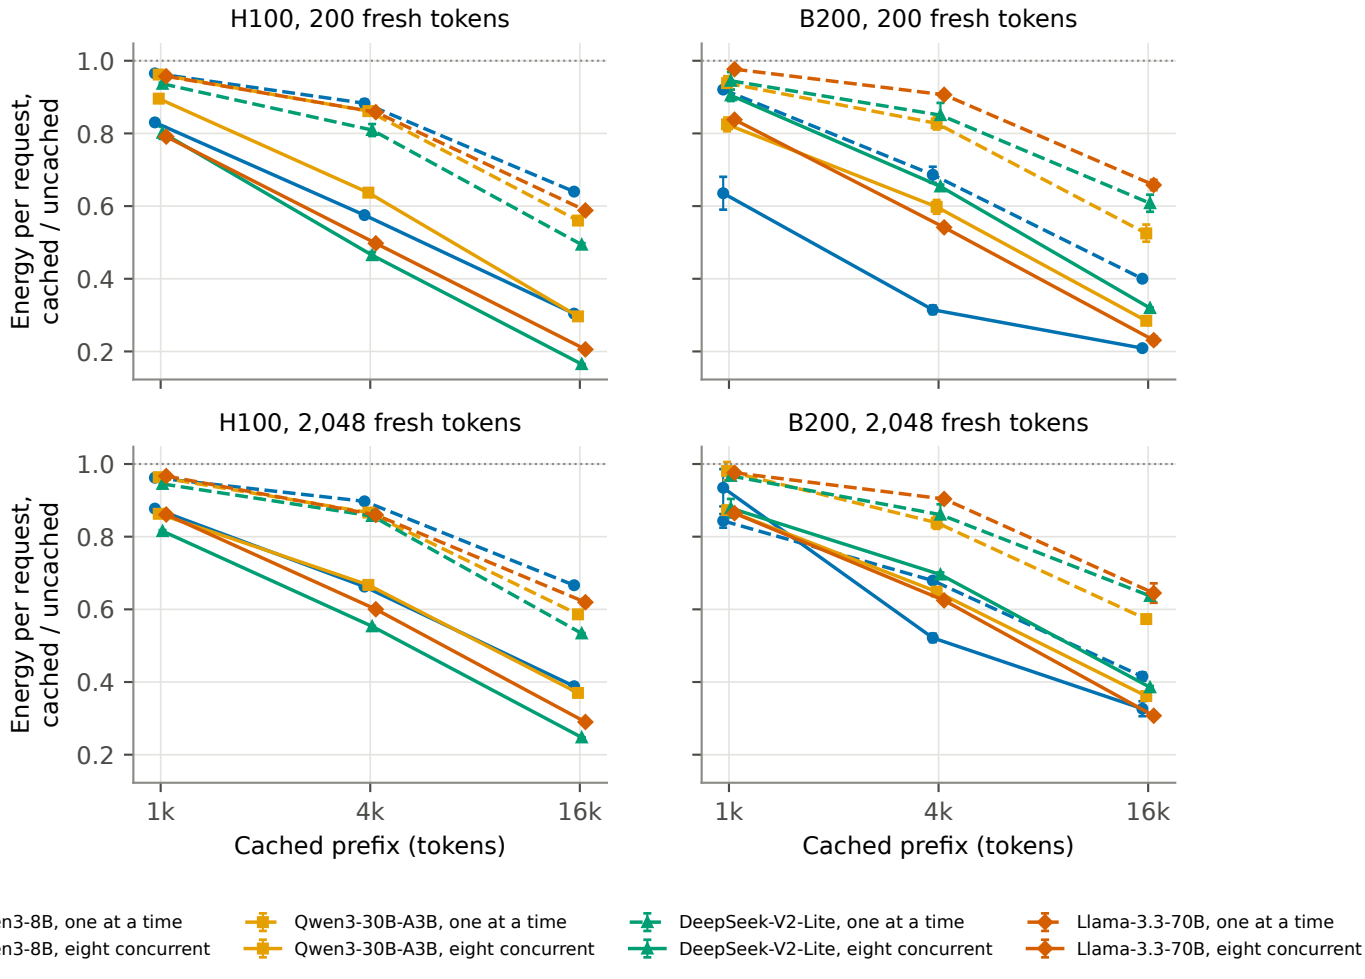
H100, 200 fresh tokens
B200, 200 fresh tokens
H100, 2,048 fresh tokens
B200, 2,048 fresh tokens
Energy per request, cached / uncached
1.0
0.8
0.6
0.4
0.2
1k
4k
16k
Cached prefix (tokens)
Qwen3-8B, one at a time
Qwen3-8B, eight concurrent
Qwen3-30B-A3B, one at a time
Qwen3-30B-A3B, eight concurrent
DeepSeek-V2-Lite, one at a time
DeepSeek-V2-Lite, eight concurrent
Llama-3.3-70B, one at a time
Llama-3.3-70B, eight concurrent
Dashed: one request at a time. Solid: eight concurrent. Bars: range over pods. 200 output tokens on both sides.